

AAST Dataset - Academic Article Survey Tables Documentation

Anonymous

1 Introduction

The Academic Article Survey Table (AAST) dataset was collected using the arXiv and Semantic Scholar APIs. It underwent a series of processing steps, including manual inspection and editing.

The processing steps are summarized as follows:

- 1) Fetching Survey Papers via the Arxiv API:
- 2) Preprocessing and Extracting Tables from Latex Files:
- 3) Creating and Partitioning the Golden Table:
- 4) Generating Descriptions for Column Headers:
- 5) Acquiring Citation Data:

2 Provided Files

We offer three files: the main file named "AAST.json", as well as two reference document detail files: "corpus_id_introduction" and "corpus_id_ori_and_sc_other".

- **AAST.json**: This is our primary file, with its content sourced from both the AAST section of the survey paper and the basic details of the references. In addition, we also provide abstracts of the references and fields generated by GPT-4-8k([OpenAI, 2023](#)) for further elucidation.
- **corpus_id_introduction.json**: Utilizing the S2ORC([Ammar et al., 2018](#)) database, this file extracts the introduction section of each reference based on its corpus_id. Both the original and the compressed versions by GPT-3.5-16k are provided for each document.
- **corpus_id_ori_and_sc_other.json**: Leveraging the S2ORC database, this file fetches other sections of the reference papers (excluding any images and tables) based on their corpus_id. Both the original and the GPT-3.5-16k compressed versions are available for each section.

3 Field Description of AAST

"<arxiv_id>", "<table_id>", "<row>", "<column>" are variable placeholders.

```
1 {  
2   "id" = "<arxiv_id>-<table_id>-<row>-<column>",  
3  
4   "arxiv_id" : "arxiv_id" ,  
5   "table_id" : "table_id (e.g. 0, 1, 1; can mappping with Golden Table)"  
6  
7 }
```

```

8   "chatGPT_field_description": "all supplementary information",
9   "chatGPT_field_c_description" : "Column Name (extraction from all
    supplementary information)",
10  "chatGPT_field_cd_description" : "Column Name + Data Type (extraction from
    all supplementary information)",
11  "chatGPT_field_cdp_description" : "Column Name + Data Type + Propose (
    extraction from all supplementary information)",
12  "chatGPT_field_cde_description" : "Column Name + Data Type + Example (
    extraction from all supplementary information)",
13  "chatGPT_field_cdpe_description" : "Column Name + Data Type + Propose +
    Example (extraction from all supplementary information)",
14
15
16  "ref_title": "The Title of the referenced literature.",
17  "ref_abstract": "Abstract of the referenced literature.",
18  "corpus_id": "ID of the reference in S2ORC of Semantic Scholar. You can use
    this corpus id to search detail information at S2ROC dataset or S2AG API.
    'https://github.com/allenai/s2orc' and 'https://www.semanticscholar.org/
    product/api'",
19
20  "ref_paper_id": "paper ID of the reference in Semantic Scholar",
21
22  "ref_year" : "Publication year of the referenced literature.",
23  "ref_authors": "Author list of the referenced literature.",
24  "ref_apa" : "The generated in-text citation in APA format.",
25  "fulltext": "Fulltext of paper",
26  "gold_header_name_and_answer_pair_tsv": "A TSV format containing table headers
    and answers: the first row represents the table headers, while the second
    row contains the answers.",
27  "gold_header_name_and_answer_pair_json": "A JSON format where the table header
    serves as the key and its corresponding answer as the value."
28
29 }

```

3.1 AAST Example

```

1
2 {
3   "id": "1904.05046v3-2-5-1",
4   "arxiv_id": "1904.05046v3",
5   "table_id": "2",
6   "paper_title": "Generalizing from a Few Examples: A Survey on Few-Shot
    Learning",
7   "table_title": "Table 5. Characteristics of embedding learning methods.",
8   "chatGPT_field_description": "Column Name: category\nData Type: String\n
    Purpose: This column indicates the type of method used for embedding
    learning, which can be task-specific, task-invariant, or hybrid. Task-
    specific methods are designed for a particular task, while task-invariant
    methods can be applied to various tasks. Hybrid methods combine elements
    of both task-specific and task-invariant approaches.\nExample: task-
    specific, task-invariant, hybrid\n\nColumn Name: method\nData Type: String
    \nPurpose: This column lists the name of the embedding learning method
    along with the author(s) and publication year. These methods are used to

```

learn representations of data in a lower-dimensional space, which can be useful for tasks such as classification, clustering, or similarity search.
 .\nExample: Method A (Author et al., Year), Method B (Author, Year), Method C (Author, Year)\n\nColumn Name: embedding function f for x test\n\nData Type: String\n\nPurpose: This column describes the function used to generate embeddings for the test data. The function can be a type of neural network (e.g., CNN, LSTM, GNN) or another method (e.g., kernel, logistic projection, adaptive CNN).\n\nExample: Neural Network A, Method D, Method E\n\nColumn Name: embedding function g for D train\n\nData Type: String\n\nPurpose: This column describes the function used to generate embeddings for the training data. The function can be the same as the one used for test data or a different one (e.g., biLSTM, another CNN).\n\nExample: the same as f, Neural Network B, Neural Network C\n\nColumn Name: similarity measure s\n\nData Type: String\n\nPurpose: This column specifies the similarity measure used to compare embeddings. The similarity measure can be a distance metric (e.g., Distance A, Distance B, Distance C) or another method (e.g., Similarity D, Similarity E, Learned Distance).\n\nExample: Similarity D, Distance A, Distance B\n\nTable: Method",

9 "chatGPT_field_c_description": "Column Name: category\n\nColumn Name: method\n\nColumn Name: embedding function f for x test\n\nColumn Name: embedding function g for D train\n\nColumn Name: similarity measure s",

10 "chatGPT_field_cd_description": "Column Name: category\n\nData Type: String\n\nColumn Name: method\n\nData Type: String\n\nColumn Name: embedding function f for x test\n\nData Type: String\n\nColumn Name: embedding function g for D train\n\nData Type: String\n\nColumn Name: similarity measure s\n\nData Type: String",

11 "chatGPT_field_cdp_description": "Column Name: category\n\nData Type: String\n\nPurpose: This column indicates the type of method used for embedding learning, which can be task-specific, task-invariant, or hybrid. Task-specific methods are designed for a particular task, while task-invariant methods can be applied to various tasks. Hybrid methods combine elements of both task-specific and task-invariant approaches.\n\nColumn Name: method\n\nData Type: String\n\nPurpose: This column lists the name of the embedding learning method along with the author(s) and publication year. These methods are used to learn representations of data in a lower-dimensional space, which can be useful for tasks such as classification, clustering, or similarity search.\n\nColumn Name: embedding function f for x test\n\nData Type: String\n\nPurpose: This column describes the function used to generate embeddings for the test data. The function can be a type of neural network (e.g., CNN, LSTM, GNN) or another method (e.g., kernel, logistic projection, adaptive CNN).\n\nColumn Name: embedding function g for D train\n\nData Type: String\n\nPurpose: This column describes the function used to generate embeddings for the training data. The function can be the same as the one used for test data or a different one (e.g., biLSTM, another CNN).\n\nColumn Name: similarity measure s\n\nData Type: String\n\nPurpose: This column specifies the similarity measure used to compare embeddings. The similarity measure can be a distance metric (e.g., Distance A, Distance B, Distance C) or another method (e.g., Similarity D, Similarity E, Learned Distance).",

12 "chatGPT_field_cde_description": "Column Name: category\n\nData Type: String\n\nExample: task-specific, task-invariant, hybrid\n\nColumn Name: method\n\nData Type: String\n\nExample: Method A (Author et al., Year), Method B (Author, Year), Method C (Author, Year)\n\nColumn Name: embedding function f for x test\n\nData Type: String\n\nExample: Neural Network A, Method D,

Method E\n\nColumn Name: embedding function g for D train\nData Type: String\nExample: the same as f, Neural Network B, Neural Network C\n\nColumn Name: similarity measure s\nData Type: String\nExample: Similarity D, Distance A, Distance B",

13 "chatGPT_field_cdpe_description": "Column Name: category\nData Type: String\nPurpose: This column indicates the type of method used for embedding learning, which can be task-specific, task-invariant, or hybrid. Task-specific methods are designed for a particular task, while task-invariant methods can be applied to various tasks. Hybrid methods combine elements of both task-specific and task-invariant approaches.\nExample: task-specific, task-invariant, hybrid\n\nColumn Name: method\nData Type: String\nPurpose: This column lists the name of the embedding learning method along with the author(s) and publication year. These methods are used to learn representations of data in a lower-dimensional space, which can be useful for tasks such as classification, clustering, or similarity search.\nExample: Method A (Author et al., Year), Method B (Author, Year), Method C (Author, Year)\n\nColumn Name: embedding function f for x test\nData Type: String\nPurpose: This column describes the function used to generate embeddings for the test data. The function can be a type of neural network (e.g., CNN, LSTM, GNN) or another method (e.g., kernel, logistic projection, adaptive CNN).\nExample: Neural Network A, Method D, Method E\n\nColumn Name: embedding function g for D train\nData Type: String\nPurpose: This column describes the function used to generate embeddings for the training data. The function can be the same as the one used for test data or a different one (e.g., biLSTM, another CNN).\nExample: the same as f, Neural Network B, Neural Network C\n\nColumn Name: similarity measure s\nData Type: String\nPurpose: This column specifies the similarity measure used to compare embeddings. The similarity measure can be a distance metric (e.g., Distance A, Distance B, Distance C) or another method (e.g., Similarity D, Similarity E, Learned Distance).\nExample: Similarity D, Distance A, Distance B",

14 "ref_title": "Matching networks for one shot learning",

15 "ref_abstract": "Learning from a few examples remains a key challenge in machine learning. Despite recent advances in important domains such as vision and language, the standard supervised deep learning paradigm does not offer a satisfactory solution for learning new concepts rapidly from little data. In this work, we employ ideas from metric learning based on deep neural features and from recent advances that augment neural networks with external memories. Our framework learns a network that maps a small labelled support set and an unlabelled example to its label, obviating the need for fine-tuning to adapt to new class types. We then define one-shot learning problems on vision (using Omniglot, ImageNet) and language tasks. Our algorithm improves one-shot accuracy on ImageNet from 87.6% to 93.2% and from 88.0% to 93.8% on Omniglot compared to competing approaches. We also demonstrate the usefulness of the same model on language modeling by introducing a one-shot task on the Penn Treebank.",

16 "corpus_id": "8909022",

17 "ref_paper_id": "be1bb4e4aa1fcf70281b4bd24d8cd31c04864bb6",

18 "ref_year": "2016",

19 "ref_authors": "Oriol Vinyals;C. Blundell;T. Lillicrap;K. Kavukcuoglu;Daan Wierstra",

20 "ref_apa": "Vinyals et al., (2016)",

21 "fulltext": "...",

```

22 "gold_header_name_and_answer_pair_tsv": "\tcategory\tembedding function f for
    x test\tembedding function g for D train\tsimilarity measure s\r\n4\ttask-
    invariant\tCNN, LSTM\tCNN, biLSTM\tcosine similarity\r\n",
23 "gold_header_name_and_answer_pair_json": [
24     {
25         "category": "task-invariant",
26         "embedding function f for x test": "CNN, LSTM",
27         "embedding function g for D train": "CNN, biLSTM",
28         "similarity measure s": "cosine similarity"
29     }
30 ]
31 },

```

4 Prompt for Semantic Compression

Given the input context limitations of LLMs, it's impractical to embed an entire document as reference context within the prompt. For this reason, we adopted the "Semantic Compression With Large Language Models" method proposed by [Gilbert et al. \(2023\)](#) and others, and applied it to condense the academic papers. Due to some introductions in the dataset exceeding 8k tokens and considering cost implications, we opted to use GPT 3.5 16K as our model for semantic compression. Furthermore, according to their research, chatGPT 3.5 also exhibits commendable semantic compression capabilities. Utilizing the section structural information provided by the S2ORC dataset, we performed semantic compression on the text of each section.

System prompt:

- 1 You are a ChatGPT LLM trained by OpenAI to compress text.
 - 2 The compression model should purely minimize the number of characters in the compressed representation, while maintaining the semantics of the original text and preserving named entities.
 - 3 The resulting compressed text does not need to be decompressed into exactly the original text, but should capture the semantics of the original text.
 - 4 The compressed text should be able to be decompressed into a text that is semantically similar to the original text, but does not need to be identical.
-

Action Prompt:

"<text>" represents the text that needs to be compressed.

- 1 [Prompt]
 - 2 Compress the following text. Return only the compressed
 - 3 text with no additional text. Text to compress:
 - 4 [Following Text]
 - 5 <text>
-

5 Field Description of corpus_id_introduction

The file "corpus_id_introduction.json" primarily provides the Original and Semantic Compression Introduction of References. The format is defined as follows:

```

1  /*
2      "corpus_id" is a variable placeholder: ": "ID of the reference in Semantic
        Scholar. You can use this corpus id to search detail information at S2ROC
        dataset or S2AG API. 'https://github.com/allenai/s2orc' and 'https://www.
        semanticscholar.org/product/api',
3  */
4  {
5      "corpus_id": {
6          "ori_introduction": "original introduction context"
7          "sc_introduction": "Semantic Compression introduction context"
8      }
9  }

```

Example

```

1  {
2      "8909022": {
3          "ori_introduction": "Introduction: \nHumans learn new concepts with very
        little supervision -e.g. a child can generalize the concept of \"
        giraffe\" from a single picture in a book -yet our best deep learning
        systems need hundreds or thousands of examples. This motivates the
        setting we are interested in: \"one-shot\" learning, which consists of
        learning a class from a single labelled example.\n\nDeep learning has
        made major advances in areas such as speech [7], vision [13] and
        language [16], but is notorious for requiring large datasets. Data
        augmentation and regularization techniques alleviate overfitting in
        low data regimes, but do not solve it. Furthermore, learning is still
        slow and based on large datasets, requiring many weight updates using
        stochastic gradient descent. This, in our view, is mostly due to the
        parametric aspect of the model, in which training examples need to be
        slowly learnt by the model into its parameters.\n\nIn contrast, many
        non-parametric models allow novel examples to be rapidly assimilated,
        whilst not suffering from catastrophic forgetting. Some models in this
        family (e.g., nearest neighbors) do not require any training but
        performance depends on the chosen metric [1]. Previous work on metric
        learning in non-parametric setups [18] has been influential on our
        model, and we aim to incorporate the best characteristics from both
        parametric and non-parametric models -namely, rapid acquisition of new
        examples while providing excellent generalisation from common
        examples.\n\nThe novelty of our work is twofold: at the modeling level
        , and at the training procedure. We propose Matching Nets (MN), a
        neural network which uses recent advances in attention and memory that
        enable rapid learning. Secondly, our training procedure is based on a
        simple machine learning principle: test and train conditions must
        match. Thus to train our network to do rapid learning, we Besides our
        contributions in defining a model and training criterion amenable for
        one-shot learning, we contribute by the definition of tasks that can
        be used to benchmark other approaches on both ImageNet and small scale
        language modeling. We hope that our results will encourage others to
        work on this challenging problem.\n\nWe organized the paper by first
        defining and explaining our model whilst linking its several
        components to related work. Then in the following section we briefly

```

```

4         elaborate on some of the related work to the task and our model. In
        Section 4 we describe both our general setup and the experiments we
        performed, demonstrating strong results on one-shot learning on a
        variety of tasks and setups.",
5     "sc_introduction": "Introduction: Humans learn new concepts with minimal
        supervision - e.g. a child generalizes the concept of \"giraffe\" from
        a single picture in a book - yet deep learning systems require
        numerous examples. This motivates \"one-shot\" learning, which
        involves learning a class from a single labeled example.\n\nDeep
        learning has made advances in speech [7], vision [13], and language
        [16], but requires large datasets. Data augmentation and
        regularization alleviate overfitting in low data regimes but do not
        solve it. Learning is slow and based on large datasets, requiring many
        weight updates using stochastic gradient descent. This is mainly due
        to the parametric aspect of the model, where training examples need to
        be slowly learned by the model into its parameters.\n\nIn contrast,
        non-parametric models allow rapid assimilation of novel examples
        without catastrophic forgetting. Some models in this family (e.g.,
        nearest neighbors) do not require training but performance depends on
        the chosen metric [1]. Previous work on metric learning in non-
        parametric setups [18] has influenced our model, and we aim to
        incorporate the best characteristics from both parametric and non-
        parametric models - rapid acquisition of new examples while providing
        excellent generalization from common examples.\n\nOur work is novel in
        two ways: at the modeling level and the training procedure. We
        propose Matching Nets (MN), a neural network that uses recent advances
        in attention and memory for rapid learning. Our training procedure is
        based on the principle that test and train conditions must match.
        Thus, to train our network for rapid learning, we Besides defining a
        model and training criterion for one-shot learning, we contribute by
        defining benchmark tasks for ImageNet and small-scale language
        modeling. We hope our results will encourage others to work on this
        challenging problem.\n\nWe organize the paper by first defining and
        explaining our model, linking its components to related work. Then, we
        briefly elaborate on related work to the task and our model. In
        Section 4, we describe our general setup and experiments,
        demonstrating strong results on one-shot learning across various tasks
        and setups."
5 }

```

6 Field Description of corpus_id_ori_and_sc_other

The file "corpus_id_ori_and_sc_other.json" primarily provides the Original and Semantic Compression Introduction of References. The format is defined as follows:

```

1  /*
2      "<corpus_id>" is a variable placeholder: ": "ID of the reference in Semantic
        Scholar. You can use this corpus id to search detail information at S2ROC
        dataset or S2AG API. 'https://github.com/allenai/s2orc' and 'https://www.
        semanticscholar.org/product/api'",
3      "<section_name>" is a variable placeholder
4  */
5  {

```



```

6     "<corpus_id>": {
7         "<section_name>": {
8             "ori_section": "original section context"
9             "sc_section": "Semantic Compression section context"
10        }, ...
11    }, ...
12 }

```

Example

```

1 {
2     "8909022": {
3         "model": {
4             "ori_section": "Our non-parametric approach to solving one-shot
                           learning is based on two components which we describe in the
                           following subsections. First, our model architecture follows
                           recent advances in neural networks augmented with memory (as
                           discussed in Section 3). Given a (small) support set  $S$ , our model
                           defines a function  $c: S \rightarrow \mathcal{Y}$  (or classifier) for each  $S$ , i.e. a mapping
                            $S \rightarrow c(S)$ . Second, we employ a training strategy which is
                           tailored for one-shot learning from the support set  $S$ .",
5             "sc_section": "Our non-parametric approach to one-shot learning is
                           based on two components: model architecture and training strategy
                           ."
6         },
7         "model architecture": {
8             "ori_section": "In recent years, many groups have investigated ways to
                           augment neural network architectures with external memories and
                           other components that make them more \"computer-like\". We draw
                           inspiration from models such as sequence to sequence (seq2seq)
                           with attention [2], memory networks [29] and pointer networks
                           [27].\n\nIn all these models, a neural attention mechanism, often
                           fully differentiable, is defined to access (or read) a memory
                           matrix which stores useful information to solve the task at hand.
                           Typical uses of this include machine translation, speech
                           recognition, or question answering. More generally, these
                           architectures model  $P(B|A)$  where  $A$  and/or  $B$  can be a sequence (
                           like in seq2seq models), or, more interestingly for us, a set
                           [26].\n\nOur contribution is to cast the problem of one-shot
                           learning within the set-to-set framework [26].\n\nThe key point is
                           that when trained, Matching Networks are able to produce sensible
                           test labels for unobserved classes without any changes to the
                           network. More precisely, we wish to map from a (small) support set
                           of  $k$  examples of image-label pairs  $S = \{(x_i, y_i)\}_{i=1}^k$  to a
                           classifier  $c: S \rightarrow \mathcal{Y}$  ( $x$ ) which, given a test example  $x$ , defines a
                           probability distribution over outputs  $\hat{\mathcal{Y}}$ . We define the
                           mapping  $S \rightarrow c(S)(x)$  to be  $P(\hat{\mathcal{Y}}|x, S)$  where  $P$  is
                           parameterised by a neural network. Thus, when given a new support
                           set of examples  $S$  from which to one-shot learn, we simply use the
                           parametric neural network defined by  $P$  to make predictions about
                           the appropriate label  $\hat{\mathcal{Y}}$  for each test example:  $P(\hat{\mathcal{Y}}|x, S)$ .
                           In general, our predicted output class for a given input
                           unseen example  $x$  and a support set  $S$  becomes  $\arg \max_{\hat{\mathcal{Y}}} P(\hat{\mathcal{Y}}|x, S)$ .

```


Our model in its simplest form computes \hat{y} as follows: $\hat{y} = \sum_{i=1}^k a(x, x_i) y_i$ (1) where x_i, y_i are the samples and labels from the support set $S = \{(x_i, y_i)\}_{i=1}^k$, and a is an attention mechanism which we discuss below. Note that eq. 1 essentially describes the output for a new class as a linear combination of the labels in the support set. Where the attention mechanism a is a kernel on $X \times X$, then (1) is akin to a kernel density estimator. Where the attention mechanism is zero for the b furthest x_i from x according to some distance metric and an appropriate constant otherwise, then (1) is equivalent to ' $k - b$ '-nearest neighbours (although this requires an extension to the attention mechanism that we describe in Section 2.1.2). Thus (1) subsumes ...",

"sc_section": "In recent years, groups have explored ways to enhance neural network architectures with external memories and components to make them more \"computer-like\". We draw inspiration from models like seq2seq with attention [2], memory networks [29], and pointer networks [27]. In these models, a neural attention mechanism is used to access a memory matrix that stores useful information for solving tasks. This includes machine translation, speech recognition, and question answering. These architectures model $P(B|A)$, where A and/or B can be a sequence or a set [26]. Our contribution is to approach one-shot learning within the set-to-set framework [26]. The key point is that Matching Networks can produce sensible test labels for unseen classes without modifying the network. Specifically, we aim to map a support set $S = \{(x_i, y_i)\}$ of k image-label pairs to a classifier $c_S(x)$ that defines a probability distribution over outputs \hat{y} for a test example x . The mapping $S \rightarrow c_S(x)$ is parameterized by a neural network P . Thus, when given a new support set S for one-shot learning, we use the neural network P to predict the appropriate label \hat{y} for each test example x : $P(\hat{y}|x, S)$. In general, our predicted output class for an unseen input example x and support set S is $\arg \max_y P(y|x, S)$. Our model computes \hat{y} as follows: $\hat{y} = \sum a(x, x_i) * y_i$ (1) where x_i, y_i are the samples and labels from the support set $S = \{(x_i, y_i)\}$ and a is an attention mechanism. Note that eq. 1 describes the output for a new class as a linear combination of the labels in the support set. If the attention mechanism a is a kernel on $X \times X$, then (1) is similar to a kernel density estimator. If the attention mechanism is zero for the b furthest x_i from x according to some distance metric and an appropriate constant otherwise, then (1) is equivalent to ' $k - b$ '-nearest neighbors. Thus, (1) encompasses both KDE and kNN methods. Another interpretation of (1) is that a acts as an attention mechanism and y_i act as memories associated with x_i . In this case, it can be seen as a specific type of associative memory where, given an input, we \"point\" to the corresponding example in the support set and retrieve its label. However, unlike other attentional memory mechanisms [2], (1) is non-parametric: as the support set size increases, so does the memory used. Therefore, the functional form defined by the classifier $c_S(x)$ is highly flexible and can easily adapt to any new support set."

},

```

11     "the attention kernel": {
12         "ori_section": "Equation 1 relies on choosing  $a(\cdot, \cdot)$ , the attention
            mechanism, which fully specifies the classifier. The simplest form
            that this takes (and which has very tight relationships with
            common attention models and kernel functions) is to use the
            softmax over the cosine distance  $c$ , i.e.,  $a(x, x_i) = \frac{e^{c(f(x), g(x_i))}}{\sum_{j=1}^k e^{c(f(x), g(x_j))}}$  with embedding functions  $f$  and  $g$ 
            being appropriate neural networks (potentially with  $f = g$ ) to
            embed  $x$  and  $x_i$ . In our experiments we shall see examples where  $f$ 
            and  $g$  are parameterised variously as deep convolutional networks
            for image tasks (as in VGG [22] or Inception [24]) or a simple
            form word embedding for language tasks (see Section 4).\n\nWe note
            that, though related to metric learning, the classifier defined
            by Equation 1 is discriminative. For a given support set  $S$  and
            sample to classify  $x$ , it is enough for  $x$  to be sufficiently aligned
            with pairs  $(x, y) \in S$  such that  $y = y$  and misaligned with the
            rest. This kind of loss is also related to methods such as
            Neighborhood Component Analysis (NCA) [18], triplet loss [9] or
            large margin nearest neighbor [28).\n\nHowever, the objective that
            we are trying to optimize is precisely aligned with multi-way,
            one-shot classification, and thus we expect it to perform better
            than its counterparts. Additionally, the loss is simple and
            differentiable so that one can find the optimal parameters in an
            \end-to-end\ fashion.",
13         "sc_section": "Equation 1 relies on choosing  $a(\cdot, \cdot)$ , the attention
            mechanism, which fully specifies the classifier. The simplest form
            :  $a(x, x_i) = \frac{e^{c(f(x), g(x_i))}}{\sum_{j=1}^k e^{c(f(x), g(x_j))}}$  with embedding
            functions  $f$  and  $g$  being appropriate neural networks (potentially
            with  $f = g$ ) to embed  $x$  and  $x_i$ . In our experiments we shall see
            examples where  $f$  and  $g$  are parameterised variously as deep
            convolutional networks for image tasks (as in VGG [22] or
            Inception [24]) or a simple form word embedding for language tasks
            (see Section 4).\n\nWe note that, though related to metric
            learning, the classifier defined by Equation 1 is discriminative.
            For a given support set  $S$  and sample to classify  $x$ , it is enough
            for  $x$  to be sufficiently aligned with pairs  $(x, y) \in S$  such that
             $y = y$  and misaligned with the rest. This kind of loss is also
            related to methods such as Neighborhood Component Analysis (NCA)
            [18], triplet loss [9] or large margin nearest neighbor [28).\n\n
            However, the objective we are trying to optimize is precisely
            aligned with multi-way, one-shot classification, and thus we
            expect it to perform better than its counterparts. Additionally,
            the loss is simple and differentiable so that one can find the
            optimal parameters in an \end-to-end\ fashion."
14     },
15     "full context embeddings": {
16         "ori_section": "The main novelty of our model lies in reinterpreting a
            well studied framework (neural networks with external memories)
            to do one-shot learning. Closely related to metric learning, the
            embedding functions  $f$  and  $g$  act as a lift to feature space  $X$  to
            achieve maximum accuracy through the classification function
            described in eq. 1.\n\nDespite the fact that the classification
            strategy is fully conditioned on the whole support set through  $P$ 
            ( $\cdot|x, S$ ), the embeddings on which we apply the cosine similarity

```

to `\\"attend\\", \\"point\\"` or simply compute the nearest neighbor are myopic in the sense that each element x_i gets embedded by $g(x_i)$ independently of other elements in the support set S . Furthermore, S should be able to modify how we embed the test image x through f . We propose embedding the elements of the set through a function which takes as input the full set S in addition to x_i , i.e. g becomes $g(x_i, S)$. Thus, as a function of the whole support set S , g can modify how to embed x_i . This could be useful when some element x_j is very close to x_i , in which case it may be beneficial to change the function with which we embed x_i - some evidence of this is discussed in Section 4. We use a bidirectional Long-Short Term Memory (LSTM) [8] to encode x_i in the context of the support set S , considered as a sequence (see appendix for a more precise definition). The second issue can be fixed via an LSTM with read-attention over the whole set S , whose inputs are equal to x : $\text{nf}(x, S) = \text{attLSTM}(f(x), g(S), K)$, where $f(x)$ are the features (e.g., derived from a CNN) which are input to the LSTM (constant at each time step). K is the fixed number of unrolling steps of the LSTM, and $g(S)$ is the set over which we attend, embedded with g . This allows for the model to potentially ignore some elements in the support set S , and adds `\\"depth\\"` to the computation of attention (see appendix for more details).",

`"sc_section"`: "The main novelty of our model is reinterpreting a well-studied framework (NNs with external memories) for one-shot learning. Related to metric learning, the embedding functions f and g lift features to space X for maximum accuracy in the classification function (eq. 1). Despite the classification strategy being conditioned on the support set through $P(.|x, S)$, the embeddings are myopic. Each element x_i is embedded independently by $g(x_i)$ regardless of other elements in S . Additionally, S can modify how we embed the test image x through f . We propose embedding the set elements using a function that takes S and x_i as input, i.e., g becomes $g(x_i, S)$. Thus, g can modify how x_i is embedded based on the support set S . This is useful when x_i is close to x_j , allowing us to change the embedding function for x_i . We use a bidirectional LSTM [8] to encode x_i in the context of S , treated as a sequence. The second issue is addressed by an LSTM with read-attention over S , where the inputs are x : $\text{nf}(x, S) = \text{attLSTM}(f(x), g(S), K)$, where $f(x)$ are the features (e.g., from a CNN) input to the LSTM (constant at each time step). K is the fixed number of unrolling steps of the LSTM, and $g(S)$ is the attended set embedded with g . This allows the model to potentially ignore some elements in S and adds `\\"depth\\"` to attention computation."

`},...`

`}`

`,...`

`}`

References

- Ammar, W., Groeneveld, D., Bhagavatula, C., Beltagy, I., Crawford, M., Downey, D., Dunkelberger, J., Elgohary, A., Feldman, S., Ha, V., Kinney, R., Kohlmeier, S., Lo, K., Murray, T., Ooi, H.-H., Peters, M., Power, J., Skjongsberg, S., Wang, L. L., Wilhelm, C., Yuan, Z., van Zuylen, M., and Etzioni, O. (2018). Construction of the literature graph in semantic scholar. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 3 (Industry Papers)*, pages 84–91, New Orleans - Louisiana. Association for Computational Linguistics.
- Gilbert, H., Sandborn, M., Schmidt, D. C., Spencer-Smith, J., and White, J. (2023). Semantic compression with large language models. *ArXiv*, abs/2304.12512.
- OpenAI (2023). Gpt-4 technical report.